

Introduction

This is part 2 of a two part 'end-of-phase' project with Flatiron School. In the previous notebook I conducted an descriptive and inferential analysis of the dataset that I will be using in this model creation.

✓ Recap

In the first notebook I answered many questions both descriptive and inferential about the data within my dataset here. I think the question that will be most apparent in this notebook is 'Which Companies Account for the highest amount of Market Capitalization?' I will later use that question first to create a failed model and then to use as the outputs for my working model

✓ Table of Contents

- [Introduction](#)
- [Recap](#)
- [Data Cleaning and Preprocessing](#)
- [Shotgun Method](#)
 - [Advanced Machine Learning: 35 Stocks](#)
- [Advanced Machine Learning: Single Stock Futures](#)
- [Creating The Pipeline](#)
- [The Final Result](#)
- [Summary](#)

```
import nbformat

def generate_toc(notebook_path):
    with open(notebook_path) as f:
        nb = nbformat.read(f, as_version=4)

    toc = []
    for cell in nb.cells:
        if cell.cell_type == 'markdown':
            lines = cell.source.split('\n')
            for line in lines:
                if line.startswith('#'):
                    header_level = line.count('#')
                    header_text = line.replace('#', '').strip()
                    toc.append((header_level, header_text))

    toc_md = ['## Table of Contents']
    for level, text in toc:
        toc_md.append(f"{' ' * (level - 1)}- [{text}](#{text.replace(' ', '-')})")

    return '\n'.join(toc_md)

notebook_path = 'stock model.ipynb'
toc_md = generate_toc(notebook_path)

# Print the generated TOC
# print(toc_md)
```

✓ Data Cleaning and Preprocessing

Lets begin!

Imports

```
import pandas as pd
import os
import tensorflow as tf
import matplotlib.pyplot as plt
```

```

import numpy as np
import seaborn as sns
from keras.models import Sequential, load_model
from keras.layers import LSTM, Dense, Dropout, BatchNormalization, Bidirectional
from keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
from keras.regularizers import l2
from keras.optimizers import Adam
from keras.utils import plot_model
from sklearn.preprocessing import MinMaxScaler, LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_squared_error, explained_variance_score, mean_absolute_error, root_mean_squared_error, mean_absol
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
import yfinance as yf
from datetime import datetime, timedelta
import warnings
from math import pi
import io
warnings.filterwarnings('ignore')

```

```
df_stocks = pd.read_csv('data/sp500_stocks.csv')
```

Im going to use the data from the S&P 500 dataset simply to get the tickers my model

```

sp500_tickers = df_stocks['Symbol'].unique()
tickers = []
for x in sp500_tickers:
    tickers.append(x)

```

Now im using yfinance (yahoo finance api) to collect the data I want to use for shotgun method model prediction

```

# Keys you're interested in from the 'info' function
keys = ['country', 'sector', 'previousClose', 'overallRisk', 'beta', 'trailingPE', 'forwardPE', 'volume', 'marketCap',
        'enterpriseValue', 'profitMargins', 'sharesOutstanding', 'bookValue',
        'priceToBook', 'trailingEps', 'forwardEps', 'pegRatio', 'enterpriseToRevenue',
        'enterpriseToEbitda', 'totalCash', 'totalCashPerShare', 'ebitda', 'totalDebt',
        'quickRatio', 'currentRatio', 'totalRevenue', 'debtToEquity', 'revenuePerShare',
        'returnOnAssets', 'returnOnEquity', 'freeCashflow', 'operatingCashflow',
        'earningsGrowth', 'revenueGrowth', 'grossMargins', 'ebitdaMargins',
        'operatingMargins', 'trailingPegRatio'
]

```

```

# Empty list to store results
data = []

```

```

# Loop through tickers and get the specific info
for ticker in tickers:
    stock = yf.Ticker(ticker)
    info = stock.info
    # Extract the values of the keys you're interested in
    row = {key: info.get(key, None) for key in keys}
    data.append(row)

```

```

# Convert the list of dictionaries into a pandas DataFrame
df = pd.DataFrame(data)

```

```

# Initialize the LabelEncoder
label_encoder = LabelEncoder()

```

```

# Apply label encoding to both columns and replace the original ones
df['country'] = label_encoder.fit_transform(df['country'])
df['sector'] = label_encoder.fit_transform(df['sector'])

```

```

# Initialize the MinMaxScaler
scaler = StandardScaler()

```

```

# Columns to scale
columns_to_scale = ['previousClose', 'overallRisk', 'beta', 'trailingPE', 'forwardPE', 'volume', 'marketCap',

```

```

'enterpriseValue', 'profitMargins', 'sharesOutstanding', 'bookValue',
'priceToBook', 'trailingEps', 'forwardEps', 'pegRatio', 'enterpriseToRevenue',
'enterpriseToEbitda', 'totalCash', 'totalCashPerShare', 'ebitda', 'totalDebt',
'quickRatio', 'currentRatio', 'totalRevenue', 'debtToEquity', 'revenuePerShare',
'returnOnAssets', 'returnOnEquity', 'freeCashflow', 'operatingCashflow',
'earningsGrowth', 'revenueGrowth', 'grossMargins', 'ebitdaMargins',
'operatingMargins', 'trailingPegRatio'
]

```

```

# Apply MinMaxScaler to the selected columns
df[columns_to_scale] = scaler.fit_transform(df[columns_to_scale])

```

```
df.dropna(inplace=True)
```

```
features = df.drop(columns=['previousClose'])
```

✓ Shotgun Method

This method is used to create a baseline accuracy from a number of models all at once

```

X = features
y = df['previousClose']

```

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.25,random_state=42)
```

```

def train_and_evaluate_models(models):
    metrics = {
        'R-squared': r2_score,
        'Mean Squared Error': mean_squared_error,
        'Mean Absolute Error': mean_absolute_error,
        'Explained Variance': explained_variance_score,
        'Root Mean Squared Error': root_mean_squared_error,
        'Mean Absolute Percentage Error': mean_absolute_percentage_error
    }

```

```
results = {}
```

```

for metric_name, metric_function in metrics.items():
    model_scores = {}
    for model in models:
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
        score = metric_function(y_test, y_pred)
        model_scores[str(model)] = score
    results[metric_name] = model_scores

```

```
return results
```

```

models = [RandomForestRegressor(), LinearRegression(), KNeighborsRegressor()]
model_metrics = train_and_evaluate_models(models=models)

```

```
fig, axes = plt.subplots(6, 1, figsize=(20, 15))
```

```
# List of metric names for easier iteration
```

```
metric_names = ['R-squared', 'Explained Variance', 'Mean Squared Error', 'Mean Absolute Error', 'Root Mean Squared Error', 'Mean Absolute Pe
```

```
for i, metric_name in enumerate(metric_names):
```

```

    ax = axes[i]
    model_scores = model_metrics[metric_name]

```

```
sns.barplot(x=list(model_scores.keys()), y=list(model_scores.values()), palette="viridis", ax=ax)
```

```
# Add labels to the bars
```

```

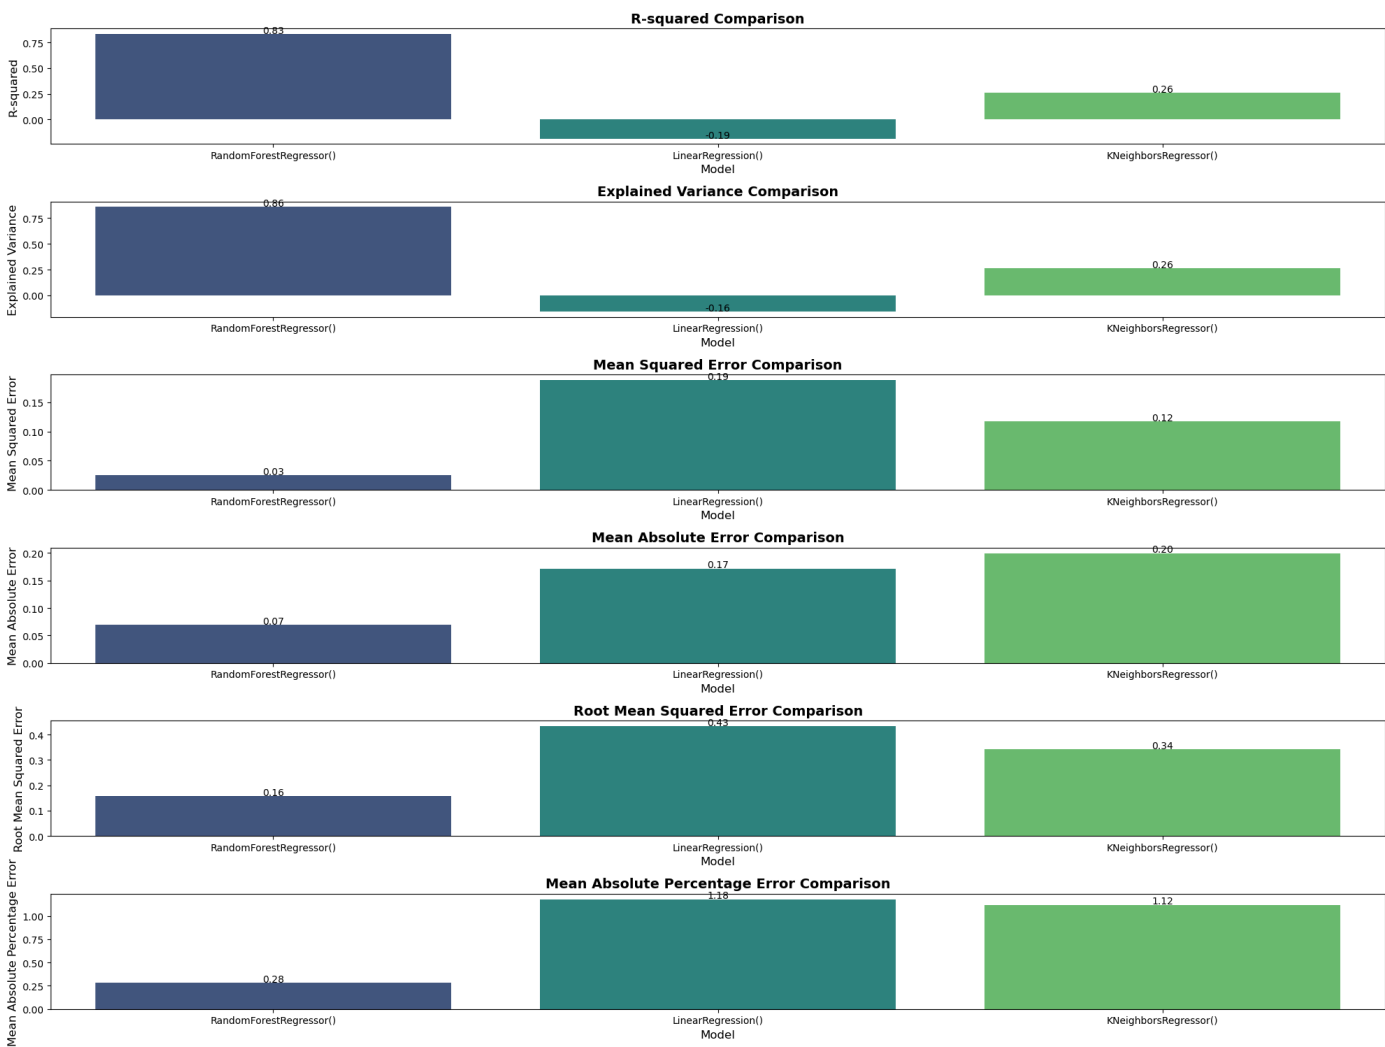
for j, score in enumerate(model_scores.values()):
    if metric_name == 'R-squared':
        ax.text(j, score + 0.01 * max(model_scores.values()), f'{score:.2f}', ha='center')
    else:
        ax.text(j, score + 0.01 * max(model_scores.values()), f'{score:.2f}', ha='center')

```

```
# Adding labels and title
```

```
ax.set_xlabel('Model', fontsize=12)
ax.set_ylabel(metric_name, fontsize=12)
ax.set_title(f'{metric_name} Comparison', fontsize=14, weight='bold')

# Adjust layout to prevent overlapping
plt.tight_layout()
# plt.savefig('savefig/RegressionModels.png')
plt.show()
```



```

# Reorder the list of metric names so that 'R-squared' and 'Explained Variance' are next to each other
metric_names = ['R-squared', 'Explained Variance', 'Mean Squared Error', 'Mean Absolute Error', 'Root Mean Squared Error', 'Mean Absolute Pe

# Prepare the data for radar chart
model_names = list(next(iter(model_metrics.values())).keys()) # Get the list of model names
num_metrics = len(metric_names)

# Initialize a dictionary to hold the scaled values
scaled_model_metrics = {metric: {} for metric in metric_names}

# Apply MinMax scaling for each metric individually across models
scaler = MinMaxScaler()

for metric in metric_names:
    # Extract the scores for this metric across all models
    scores = np.array([model_metrics[metric][model] for model in model_names]).reshape(-1, 1)

    # Scale the scores for this metric
    scaled_scores = scaler.fit_transform(scores).flatten()

    # Store the scaled scores back into the dictionary
    for i, model in enumerate(model_names):
        scaled_model_metrics[metric][model] = scaled_scores[i]

# Create angles for the radar plot to form a hexagon (6 sides)
angles = [n / float(num_metrics) * 2 * pi for n in range(num_metrics)]
angles += angles[:1] # Complete the loop

# Define the color mapping for each model with exact colors
color_mapping = {
    'RandomForestRegressor()': '#4CAF50', # Green used previously
    'KNeighborsRegressor()': 'yellow',
    'LinearRegression()': '#F44336' # Red used previously
}

# Create the radar plot with hexagonal grid
fig, ax = plt.subplots(figsize=(10, 10), subplot_kw=dict(polar=True))

# Customize the hexagonal appearance by manually setting the radial ticks and limits
ax.set_theta_offset(pi / 6) # Start the first axis at the top of the hexagon
ax.set_theta_direction(-1) # Ensure the axes go clockwise

# Plot each model's scaled metrics with their assigned colors
for model in model_names:
    scores = [scaled_model_metrics[metric][model] for metric in metric_names]
    scores += scores[:1] # Complete the loop for radar chart

    ax.plot(angles, scores, linewidth=2, linestyle='solid', label=model, color=color_mapping.get(str(model), 'blue'))
    ax.fill(angles, scores, alpha=0.25, color=color_mapping.get(str(model), 'blue'))

# Add labels for each metric at the corresponding angles
plt.xticks(angles[:-1], metric_names, color='black', size=14, weight='bold')

# Set y-ticks and y-limits for better readability (since everything is now scaled between 0 and 1)
plt.yticks([0.2, 0.4, 0.6, 0.8, 1.0], ['0.2', '0.4', '0.6', '0.8', '1.0'], color="grey", size=12)
plt.ylim(0, 1)

# Customize the grid to have a hexagonal shape
ax.grid(color='grey', linestyle='--', linewidth=0.5, alpha=0.7)
ax.spines['polar'].set_visible(False)

# Title and legend
plt.title('Stock Analysis: Model Performance Comparison', size=18, color='black', weight='bold', pad=20)


# Move the legend to the lower-right corner
plt.legend(loc='lower right', bbox_to_anchor=(1.1, -0.1), fontsize=12)

# Display the plot
plt.tight_layout()
plt.savefig('shotgun.png', transparent=True)
plt.show()

```



Stock Analysis: Model Performance Comparison



Advanced Machine Learning: 35 Stocks

This was the original concept I conceived, the plan was to train the model on 35 different stocks at once and then use that model to predict new data for a stock that wasn't in the original 35, it went just about as well as it sounds

```
# List of 35 stocks with the highest market cap in respective sector
tickers = ['AAPL', 'MSFT', 'NVDA', 'AVGO', 'ORCL', 'ADBE', 'CRM', 'AMD', 'ACN', # 9 Stocks from the Technology Sector
           'BRK-B', 'JPM', 'V', 'MA', 'BAC', 'WFC', # 6 Stocks from the Financial Services Sector
           'LLY', 'UNH', 'JNJ', 'ABBV', # 4 Stocks from the Healthcare Sector
           'GE', 'CAT', 'RTX', 'UNP', # 4 Stocks from the Industrials Sector
           'AMZN', 'TSLA', 'HD', # 3 Stocks from the Consumer Cyclical Sector
           'GOOGL', 'META', 'NFLX', # 3 Stocks from the Communication Services Sector
           'WMT', 'PG', # 2 Stocks from the Consumer Defensive Sector
           'XOM', # 1 Stock from the Energy Sector
           'PLD', # 1 Stock from the Real Estate Sector
           'NEE', # 1 Stock from the Utilities Sector
           'LIN' # 1 Stock from the Basic Materials Sector
```

```
]
```

```
# Define the end date as yesterday
end_date = datetime.strptime('2023-09-13', '%Y-%m-%d')
```



```
X_train_resaped.shape, y_train_resaped.shape
```

```
((92085, 1, 60), (92085,))
```

```
model = Sequential()
model.add(LSTM(100, return_sequences=True, input_shape=(1, time_steps)))
model.add(Dropout(0.2))
model.add(Bidirectional(LSTM(100, return_sequences=True)))
model.add(Dropout(0.2))
model.add(Bidirectional(LSTM(100, return_sequences=True)))
model.add(Dropout(0.2))
model.add(LSTM(100))
model.add(Dropout(0.2))
model.add(Dense(1))
```

```
optimizer = Adam(learning_rate=0.000001)
model.compile(optimizer=optimizer, loss='mean_squared_error')
```

```
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 1, 100)	64400
dropout (Dropout)	(None, 1, 100)	0
bidirectional (Bidirectional)	(None, 1, 200)	160800
dropout_1 (Dropout)	(None, 1, 200)	0
bidirectional_1 (Bidirectional)	(None, 1, 200)	240800
dropout_2 (Dropout)	(None, 1, 200)	0
lstm_3 (LSTM)	(None, 100)	120400
dropout_3 (Dropout)	(None, 100)	0
dense (Dense)	(None, 1)	101
Total params: 586,501		
Trainable params: 586,501		
Non-trainable params: 0		

i commented the fitting of my model as the best model has been saved and I wont accidentally spend the time retraining it

```
# # Early stopping callback
# early_stop = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)

# # Train the model with the checkpoint
# history = model.fit(X_train_resaped, y_train_resaped, epochs=50, batch_size=10, validation_split=0.2, shuffle=False, callbacks=[early_st

# model.save('best_model.h5')

model= load_model('best_model.h5')

def stock_prediction(test_ticker):
    # Load test stock data
    test_stock = pd.DataFrame()
    df = yf.download(test_ticker, start="2014-01-01", end=end_date_str)
    df['Ticker'] = test_ticker
    test_stock = pd.concat([test_stock, df])

    # Select and pivot the test data
    test_stock = test_stock[['Ticker', 'Adj Close']]
    test_stock = test_stock.pivot_table(index=test_stock.index, columns='Ticker', values='Adj Close')

    # Scaling the test data
    scaler = MinMaxScaler()
```

```
scaled_test_data = pd.DataFrame(scaler.fit_transform(test_stock), columns=test_stock.columns, index=test_stock.index)

# Create sequences
X_test, y_test = [], []
for i in range(len(scaled_test_data) - time_step):
    X_test.append(scaled_test_data.iloc[i:i + time_step].values)
    y_test.append(scaled_test_data.iloc[i + time_step].values)

X_test = np.array(X_test)
y_test = np.array(y_test)


# Reshape the test data
X_test_reshaped = X_test.reshape(-1, 1, time_step)
y_test_reshaped = y_test.reshape(-1)

# Make predictions
predictions = model.predict(X_test_reshaped)

# Plot the predictions vs actual stock prices starting from the 2000th sample
plt.plot(y_test_reshaped[2455:], label='Actual Stock Prices')
plt.plot(predictions[2455:], label='Predicted Stock Prices', linestyle='--')
plt.legend()
plt.title(f'Stock Price Prediction vs Actual for {test_ticker}')
plt.show()
```

```
stock_prediction('AMZN')
```

```
→ [*****100%*****] 1 of 1 completed
83/83 [=====] - 2s 2ms/step
```



as you can see, the results were not good

✓ Advanced Machine Learning: Single Stock Futures

```
# Define the end date as yesterday
end_date = datetime.now() - timedelta(1)
end_date_str = end_date.strftime('%Y-%m-%d')

# Define the start date as a year before yesterday
start_date = end_date - timedelta(1825)
start_date_str = start_date.strftime('%Y-%m-%d')

# Download stock data
stock_data = yf.download('WM', start=start_date_str, end=end_date_str)
```

```
# Scaling
scaler = MinMaxScaler(feature_range=(-1, 1))
scaled_data = scaler.fit_transform(stock_data['Close'].values.reshape(-1,1))
```

100%*****] 1 of 1 completed

```
time_step = 7
```

```
# Define the model
model = Sequential()
model.add(LSTM(units=64, return_sequences=True, input_shape=(time_step, 1)))
model.add(Bidirectional(LSTM(units=64)))
model.add(Dense(units=64))
model.add(Dropout(0.5))
model.add(Dense(units=1))
```

```
# Compile the model
optimizer = Adam(learning_rate=0.001)
model.compile(optimizer=optimizer, loss='mean_squared_error')
```

```
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
lstm_4 (LSTM)	(None, 7, 64)	16896
bidirectional_2 (Bidirectional)	(None, 128)	66048
dense_1 (Dense)	(None, 64)	8256
dropout_4 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 1)	65
=====		
Total params: 91,265		
Trainable params: 91,265		
Non-trainable params: 0		

```
# Step 1: Predict today's data
todays_data = scaled_data[-time_step:].reshape(1, time_step, 1)
todays_prediction = model.predict(todays_data)
todays_prediction_inversed = scaler.inverse_transform(todays_prediction)
todays_prediction_inversed
```

1/1 [=====] - 1s 1s/step
array([[150.05959]], dtype=float32)

```
# Step 2: Append today's prediction to the dataset
new_data_point = np.append(stock_data['Close'].values, todays_prediction_inversed)
```

```
# Recreate scaled_data based on the updated dataset
updated_scaled_data = scaler.fit_transform(new_data_point.reshape(-1, 1))
```

```
# Recreate X and y with the updated dataset
X_new, y_new = [], []
for i in range(len(updated_scaled_data) - time_step - 1):
    X_new.append(updated_scaled_data[i:(i + time_step), 0]) # Add time_step-length sequences to X
    y_new.append(updated_scaled_data[i + time_step, 0]) # Add the next value to y

X_new = np.array(X_new)
y_new = np.array(y_new)
X_new = np.reshape(X_new, (X_new.shape[0], X_new.shape[1], 1))
```

```
# Early stopping callback
early_stop = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)
```

```
# Learning rate reducer callback
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=5, min_lr=1e-10, verbose=1)
```

```
# Train the model with the checkpoint
history = model.fit(X_new, y_new, epochs=100, batch_size=10, validation_split=0.2, shuffle=False, callbacks=[early_stop, reduce_lr])
```

```
Epoch 1/100
100/100 [=====] - 1s 9ms/step - loss: 0.0031 - val_loss: 0.0017 - lr: 1.0000e-06
Epoch 2/100
100/100 [=====] - 1s 8ms/step - loss: 0.0030 - val_loss: 0.0017 - lr: 1.0000e-06
Epoch 3/100
100/100 [=====] - 1s 9ms/step - loss: 0.0036 - val_loss: 0.0017 - lr: 1.0000e-06
Epoch 4/100
100/100 [=====] - 1s 8ms/step - loss: 0.0033 - val_loss: 0.0016 - lr: 1.0000e-06
Epoch 5/100
100/100 [=====] - 1s 8ms/step - loss: 0.0030 - val_loss: 0.0016 - lr: 1.0000e-06
Epoch 6/100
98/100 [=====>.] - ETA: 0s - loss: 0.0033
Epoch 6: ReduceLROnPlateau reducing learning rate to 1.000000111620805e-07.
100/100 [=====] - 1s 8ms/step - loss: 0.0033 - val_loss: 0.0016 - lr: 1.0000e-07
Epoch 7/100
100/100 [=====] - 1s 8ms/step - loss: 0.0029 - val_loss: 0.0016 - lr: 1.0000e-07
Epoch 8/100
100/100 [=====] - 1s 14ms/step - loss: 0.0031 - val_loss: 0.0016 - lr: 1.0000e-07
Epoch 9/100
100/100 [=====] - 1s 8ms/step - loss: 0.0033 - val_loss: 0.0016 - lr: 1.0000e-07
Epoch 10/100
100/100 [=====] - 1s 8ms/step - loss: 0.0032 - val_loss: 0.0016 - lr: 1.0000e-07
Epoch 11/100
98/100 [=====>.] - ETA: 0s - loss: 0.0033
Epoch 11: ReduceLROnPlateau reducing learning rate to 1.000000082740371e-08.
100/100 [=====] - 1s 8ms/step - loss: 0.0032 - val_loss: 0.0016 - lr: 1.0000e-07
Epoch 12/100
100/100 [=====] - 1s 8ms/step - loss: 0.0036 - val_loss: 0.0016 - lr: 1.0000e-08
Epoch 13/100
100/100 [=====] - 1s 8ms/step - loss: 0.0031 - val_loss: 0.0016 - lr: 1.0000e-08
Epoch 14/100
100/100 [=====] - 1s 8ms/step - loss: 0.0036 - val_loss: 0.0016 - lr: 1.0000e-08
Epoch 15/100
100/100 [=====] - 1s 7ms/step - loss: 0.0032 - val_loss: 0.0016 - lr: 1.0000e-08
Epoch 16/100
98/100 [=====>.] - ETA: 0s - loss: 0.0035
Epoch 16: ReduceLROnPlateau reducing learning rate to 1.000000082740371e-09.
100/100 [=====] - 1s 8ms/step - loss: 0.0034 - val_loss: 0.0016 - lr: 1.0000e-08
Epoch 17/100
100/100 [=====] - 1s 8ms/step - loss: 0.0031 - val_loss: 0.0016 - lr: 1.0000e-09
Epoch 18/100
100/100 [=====] - 1s 7ms/step - loss: 0.0033 - val_loss: 0.0016 - lr: 1.0000e-09
Epoch 19/100
100/100 [=====] - 1s 8ms/step - loss: 0.0034 - val_loss: 0.0016 - lr: 1.0000e-09
Epoch 20/100
100/100 [=====] - 1s 8ms/step - loss: 0.0032 - val_loss: 0.0016 - lr: 1.0000e-09
Epoch 21/100
98/100 [=====>.] - ETA: 0s - loss: 0.0034
Epoch 21: ReduceLROnPlateau reducing learning rate to 1.000000082740371e-10.
100/100 [=====] - 1s 8ms/step - loss: 0.0033 - val_loss: 0.0016 - lr: 1.0000e-09
Epoch 22/100
100/100 [=====] - 1s 8ms/step - loss: 0.0031 - val_loss: 0.0016 - lr: 1.0000e-10
Epoch 23/100
100/100 [=====] - 1s 8ms/step - loss: 0.0035 - val_loss: 0.0016 - lr: 1.0000e-10
Epoch 24/100
100/100 [=====] - 1s 8ms/step - loss: 0.0030 - val_loss: 0.0016 - lr: 1.0000e-10
Epoch 25/100
100/100 [=====] - 1s 7ms/step - loss: 0.0034 - val_loss: 0.0016 - lr: 1.0000e-10
```

```
# Plot the training loss and validation loss with stock analysis theme
plt.figure(figsize=(10, 6))
```

```
# Plot training loss in green (typically for "positive" improvement)
plt.plot(history.history['loss'], label='Training Loss', color='#4CAF50', linewidth=2)
```

```
# Plot validation loss in red (for validation, possibly showing "issues")
plt.plot(history.history['val_loss'], label='Validation Loss', color='#F44336', linewidth=2)
```

```
# Set title and labels
plt.title('Training and Validation Loss (Stock Analysis Theme)', fontsize=18, weight='bold', pad=20)
plt.xlabel('Epochs', fontsize=14, weight='bold')
plt.ylabel('Loss', fontsize=14, weight='bold')
```

```
# Set y-axis limits from 0 to 0.05 as requested
plt.ylim(0, 0.05)
```

```
# Add grid lines for better readability
plt.grid(color='grey', linestyle='--', linewidth=0.5, alpha=0.7)
```

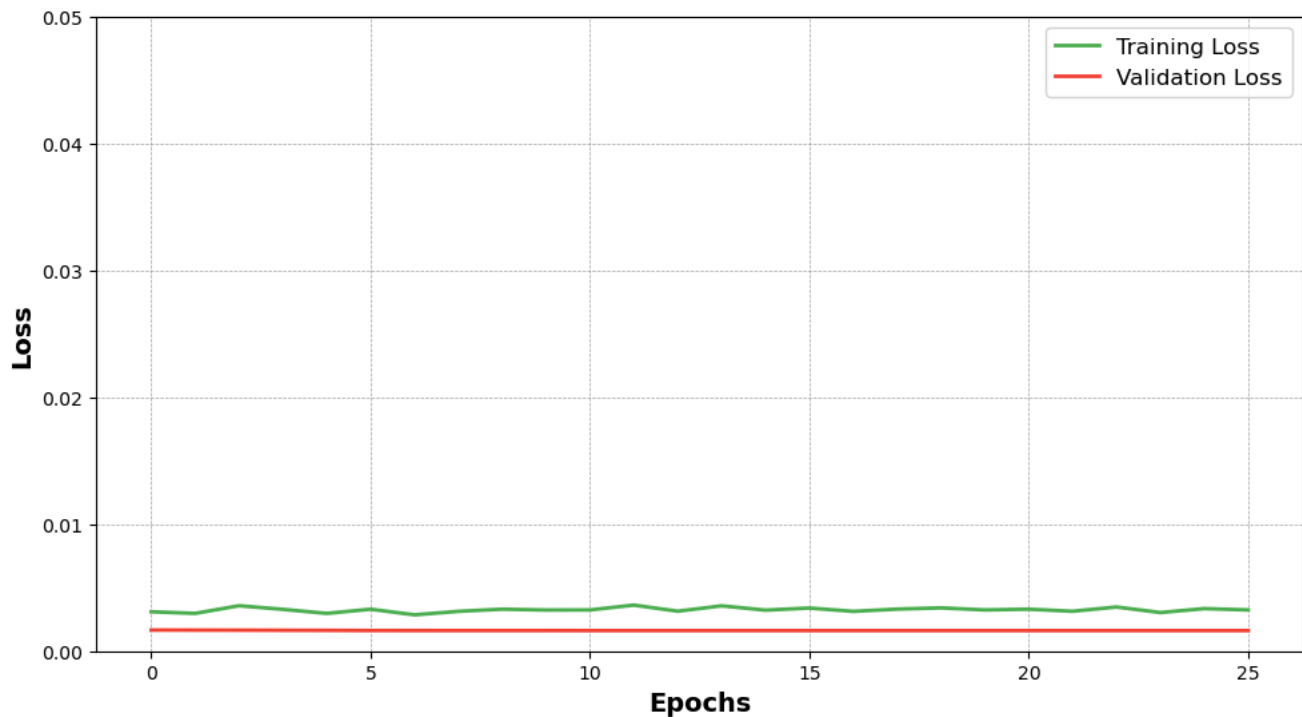
```
# Customize the legend and display it in the lower right corner
plt.legend(loc='upper right', fontsize=12)

# Clean and polished layout
plt.tight_layout()

# Show the plot
plt.savefig('Loss.png', transparent= True)
plt.show()
```



Training and Validation Loss (Stock Analysis Theme)



```
# Make new predictions after retraining
predictions = model.predict(X_new)
predictions = scaler.inverse_transform(predictions)
print(f'Predicted Close: {predictions[-1:]})')
print(f"Actual Close: {stock_data['Close'][-1:].values}")
```

```
40/40 [=====] - 0s 7ms/step
Predicted Close: [[207.55775]]
Actual Close: [207.63999939]
```

```
# Collect data for the plot
original_data = stock_data['Close'].values
predicted_data = np.empty_like(original_data)
predicted_data[:] = np.nan
predicted_data[-len(predictions):] = predictions.reshape(-1)
```

```
# Create the plot
fig, ax = plt.subplots()
ax.plot(original_data, label='Original Data')
ax.plot(predicted_data, label='Predicted Data')
ax.legend()
```

```
plt.title(f"Stock Price Prediction for Training Stock (WM)")
plt.show()
```



Stock Price Prediction for Training Stock (WM)



```
# Initialize the data with the last `time_step` days from the initial_data
initial_data = scaled_data[-7:]
current_window = initial_data[-time_step:].copy() # Assuming `initial_data` is already defined
all_predictions = []

for i in range(30):
    # Reshape current_window to match the model input shape (assuming a 3D input for LSTM)
    X_input = current_window.reshape((1, time_step, current_window.shape[1]))

    # Generate multiple predictions (ensemble) for the current day
    ensemble_predictions = []
    for _ in range(15):
        predicted_value = model.predict(X_input)[0][0] # Assuming a single output per prediction
        ensemble_predictions.append(predicted_value)

    # Calculate the mean of the ensemble predictions
    predicted_value = np.mean(ensemble_predictions)

    # Append the prediction to the list of predictions
    all_predictions.append(predicted_value)

    # Update the current window with the new predicted value
    # We assume here that we're predicting a single feature, like the "Close" price
    predicted_value_as_array = np.array([[predicted_value]])
    current_window = np.vstack([current_window[1:], predicted_value_as_array])
```



```
1/1 [=====] - ETA: 0s1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 43ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 43ms/step
```

```

1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 84ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 64ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step

```

```

# Convert predictions back to original scale (if using a scaler)
predicted_prices = scaler.inverse_transform(np.array(all_predictions).reshape(-1, 1))

# Optionally: Plot or use these predictions as needed
print("Predicted Prices for the next 14 days (using ensemble averaging):", predicted_prices)

```

```

↗ Predicted Prices for the next 14 days (using ensemble averaging): [[207.48682]
[207.3095 ]
[207.07431]
[206.83226]
[206.71588]
[206.52821]
[206.34254]
[206.16513]
[205.98631]
[205.80681]
[205.62862]
[205.45323]
[205.27734]
[205.10245]
[204.92851]
[204.75526]
[204.58284]
[204.41116]
[204.24028]
[204.07016]
[203.9008 ]
[203.73218]
[203.5643 ]
[203.39711]
[203.23068]
[203.06493]
[202.8999 ]
[202.73557]
[202.57193]
[202.40897]]

```

```

# Number of days to zoom in on
zoom_days = 730

# Create a timeline for the final `zoom_days` of the actual data and original predictions
days = np.arange(len(original_data))[-zoom_days:] # Last `zoom_days` days for the actual data
predicted_days = np.arange(len(original_data))[-zoom_days:] # Adjust for the predicted data's timeline
future_days = np.arange(len(original_data), len(original_data) + len(predicted_prices)) # Future timeline

plt.figure(figsize=(10, 6))

# Plot last `zoom_days` of actual data in blue
plt.plot(days, original_data[-zoom_days:], label='Actual Data', color='blue')

# Plot last `zoom_days` of original predictions in orange

```

```

plt.plot(predicted_days, predicted_data[-zoom_days:], label='Predicted Data', color='orange')

# Plot future predictions in green after the dotted line
plt.plot(future_days, predicted_prices, label='Future Predicted Data', color='green')

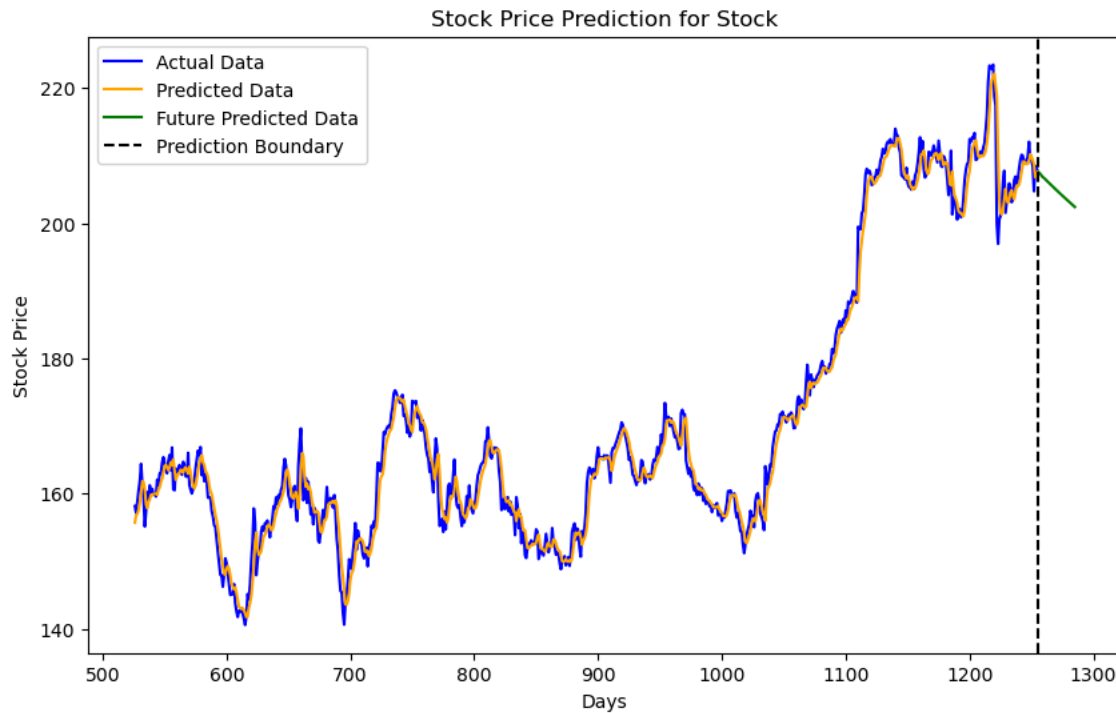
# Add a vertical dotted line at the point where the future predictions start
plt.axvline(x=len(original_data) - 1, color='black', linestyle='--', label='Prediction Boundary')

# Add labels and title
plt.xlabel('Days')
plt.ylabel('Stock Price')
plt.title('Stock Price Prediction for Stock')

# Show legend
plt.legend()

# Show plot
plt.show()

```



```

noise_factor = 0.02 # Adjust this factor to control the level of noise
initial_data = scaled_data[-7:]
current_window = initial_data[-time_step:].copy() # Assuming `initial_data` is already defined
noisy_predictions = []

```

```

for i in range(30):
    X_input = current_window.reshape((1, time_step, current_window.shape[1]))

    ensemble_predictions = []
    for _ in range(15):
        predicted_value = model.predict(X_input)[0][0]
        ensemble_predictions.append(predicted_value)

    # Calculate the mean of the ensemble predictions
    predicted_value = np.mean(ensemble_predictions)

    # Add random noise to introduce variance
    predicted_value += noise_factor * np.random.randn()

    noisy_predictions.append(predicted_value)

    predicted_value_as_array = np.array([[predicted_value]])
    current_window = np.vstack([current_window[1:], predicted_value_as_array])

```



```

1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 24ms/step

```



```
# Convert predictions back to original scale (if using a scaler)
predicted_prices = scaler.inverse_transform(np.array(noisy_predictions).reshape(-1, 1))

# Optionally: Plot or use these predictions as needed
print("Predicted Prices for the next 14 days (using ensemble averaging):", predicted_prices)
```

<https://colab.research.google.com/drive/1Wlr6fvXlahgXY1jt7xWZuZX5rmmxI5By#printMode=true>

```

[201.45977568]
[200.93608523]
[198.72822307]
[198.64837342]
[197.6346516 ]
[197.85284041]
[199.90071274]]

# Number of days to zoom in on
zoom_days = 185

# Create a timeline for the final `zoom_days` of the actual data and original predictions
days = np.arange(len(original_data))[-zoom_days:] # Last `zoom_days` days for the actual data
predicted_days = np.arange(len(original_data))[-zoom_days:] # Adjust for the predicted data's timeline
future_days = np.arange(len(original_data), len(original_data) + len(predicted_prices)) # Future timeline

plt.figure(figsize=(10, 6))

# Plot last `zoom_days` of actual data in blue
plt.plot(days, original_data[-zoom_days:], label='Actual Data', color='blue')

# Plot last `zoom_days` of original predictions in orange
plt.plot(predicted_days, predicted_data[-zoom_days:], label='Predicted Data', color='orange')

# Plot future predictions in green after the dotted line
plt.plot(future_days, predicted_prices, label='Future Predicted Data', color='green')


# Add a vertical dotted line at the point where the future predictions start
plt.axvline(x=len(original_data) - 1, color='black', linestyle='--', label='Prediction Boundary')

# Add labels and title
plt.xlabel('Days')
plt.ylabel('Stock Price')
plt.title('Stock Price Prediction for Stock')

# Show legend
plt.legend()

# Show plot
plt.show()

```



This process on the otherhand came out REALLY WELL!!

✓ Creating The Pipeline

now that I know my model works and it is performing successfully, i combined all the steps into 1 function

```
def pipeline(ticker):
    # Define the end date as yesterday
    end_date = datetime.now() - timedelta(1)
    end_date_str = end_date.strftime('%Y-%m-%d')

    # Define the start date as 5 years before yesterday
    start_date = end_date - timedelta(1825)
    start_date_str = start_date.strftime('%Y-%m-%d')

    # Define the directory path
    directory = f'Stock Graphs/{ticker}'

    # Check if the directory exists, and create it if it doesn't
    if not os.path.exists(directory):
        os.makedirs(directory)

    # Download stock data
    stock_data = yf.download(ticker, start=start_date_str, end=end_date_str)

    # Scaling
    scaler = MinMaxScaler(feature_range=(-1, 1))
    scaled_data = scaler.fit_transform(stock_data['Close'].values.reshape(-1,1))

    time_step = 7

    # Define the model
    model = Sequential()
    model.add(LSTM(units=64, return_sequences=True, input_shape=(time_step, 1)))
    model.add(Bidirectional(LSTM(units=64)))
    model.add(Dense(units=64))
    model.add(Dropout(0.5))
    model.add(Dense(units=1))

    # Compile the model
    optimizer = Adam(learning_rate=0.001)
    model.compile(optimizer=optimizer, loss='mean_squared_error')

    # Step 1: Predict today's data
    todays_data = scaled_data[-time_step:].reshape(1, time_step, 1)
    todays_prediction = model.predict(todays_data)
    todays_prediction_inversed = scaler.inverse_transform(todays_prediction)
    todays_prediction_inversed

    # Step 2: Append today's prediction to the dataset
    new_data_point = np.append(stock_data['Close'].values, todays_prediction_inversed)

    # Recreate scaled_data based on the updated dataset
    updated_scaled_data = scaler.fit_transform(new_data_point.reshape(-1, 1))

    # Recreate X and y with the updated dataset
    X_new, y_new = [], []
    for i in range(len(updated_scaled_data) - time_step - 1):
        X_new.append(updated_scaled_data[i:(i + time_step), 0]) # Add time_step-length sequences to X
        y_new.append(updated_scaled_data[i + time_step, 0])      # Add the next value to y

    X_new = np.array(X_new)
    y_new = np.array(y_new)
    X_new = np.reshape(X_new, (X_new.shape[0], X_new.shape[1], 1))

    # Early stopping callback
    early_stop = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)

    # Learning rate reducer callback
    reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=5, min_lr=1e-10, verbose=1)

    # Train the model with the checkpoint
    model.fit(X_new, y_new, epochs=100, batch_size=10, validation_split=0.2, shuffle=False, callbacks=[early_stop, reduce_lr])

    # Make new predictions after retraining
    predictions = model.predict(X_new)
```

```

predictions = scaler.inverse_transform(predictions)
print(f'Predicted Close: {predictions[-1:]})')
print(f'Actual Close: {stock_data['Close'][-1:].values}")

# Collect data for the plot
original_data = stock_data['Close'].values
predicted_data = np.empty_like(original_data)
predicted_data[:] = np.nan
predicted_data[-len(predictions):] = predictions.reshape(-1)

# Historical data plot with prediction
fig, ax = plt.subplots(figsize=(10, 6))
ax.plot(original_data, label='Original Data', color='blue', linewidth=2)
ax.plot(predicted_data, label='Predicted Data', color='orange', linewidth=2)

# Add labels and title with bold, larger fonts
ax.set_title(f'Stock Price Prediction for {ticker}', fontsize=18, weight='bold', pad=20)
ax.set_xlabel('Days', fontsize=14, weight='bold')
ax.set_ylabel('Stock Price', fontsize=14, weight='bold')

# Add a grid and legend
ax.grid(color='grey', linestyle='--', linewidth=0.5, alpha=0.7)
ax.legend(loc='lower right', fontsize=12)

plt.savefig(f'{directory}/historical.png')
plt.close()

# Initialize the data with the last `time_step` days from the initial_data
initial_data = scaled_data[-7:]
current_window = initial_data[-time_step:].copy() # Assuming `initial_data` is already defined
all_predictions = []

for i in range(30):
    # Reshape current_window to match the model input shape (assuming a 3D input for LSTM)
    X_input = current_window.reshape((1, time_step, current_window.shape[1]))

    # Generate multiple predictions (ensemble) for the current day
    ensemble_predictions = []
    for _ in range(15):
        predicted_value = model.predict(X_input)[0][0] # Assuming a single output per prediction
        ensemble_predictions.append(predicted_value)

    # Calculate the mean of the ensemble predictions
    predicted_value = np.mean(ensemble_predictions)

    # Append the prediction to the list of predictions
    all_predictions.append(predicted_value)

    # Update the current window with the new predicted value
    # We assume here that we're predicting a single feature, like the "Close" price
    predicted_value_as_array = np.array([[predicted_value]])
    current_window = np.vstack([current_window[1:], predicted_value_as_array])

# Convert predictions back to original scale (if using a scaler)
predicted_prices = scaler.inverse_transform(np.array(all_predictions).reshape(-1, 1))

# Number of days to zoom in on
zoom_days = 730

# Create a timeline for the final `zoom_days` of the actual data and original predictions
days = np.arange(len(original_data))[-zoom_days:] # Last `zoom_days` days for the actual data
predicted_days = np.arange(len(original_data))[-zoom_days:] # Adjust for the predicted data's timeline
future_days = np.arange(len(original_data), len(original_data) + len(predicted_prices)) # Future timeline

# Future predictions plot
plt.figure(figsize=(10, 6))
plt.plot(days, original_data[-zoom_days:], label='Actual Data', color='blue', linewidth=2)
plt.plot(predicted_days, predicted_data[-zoom_days:], label='Predicted Data', color='orange', linewidth=2)
plt.plot(future_days, predicted_prices, label='Future Predicted Data', color='#4CAF50', linewidth=2) # Green color for future prediction

# Add a vertical dotted line at the point where the future predictions start
plt.axvline(x=len(original_data) - 1, color='black', linestyle='--', label='Prediction Boundary')

# Add labels and title with bold fonts
plt.title(f'Stock Price Prediction for {ticker}', fontsize=18, weight='bold', pad=20)
plt.xlabel('Days', fontsize=14, weight='bold')

```

```

plt.ylabel('Stock Price', fontsize=14, weight='bold')

# Add grid and legend
plt.grid(color='grey', linestyle='--', linewidth=0.5, alpha=0.7)
plt.legend(loc='lower right', fontsize=12)

plt.savefig(f'{directory}/future.png')
plt.close()

noise_factor = 0.02 # Adjust this factor to control the level of noise
initial_data = scaled_data[-7:]
current_window = initial_data[-time_step:].copy() # Assuming `initial_data` is already defined
noisy_predictions = []

for i in range(30):
    X_input = current_window.reshape((1, time_step, current_window.shape[1]))

    ensemble_predictions = []
    for _ in range(15):
        predicted_value = model.predict(X_input)[0][0]
        ensemble_predictions.append(predicted_value)

    # Calculate the mean of the ensemble predictions
    predicted_value = np.mean(ensemble_predictions)

    # Add random noise to introduce variance
    predicted_value += noise_factor * np.random.randn()

    noisy_predictions.append(predicted_value)

    predicted_value_as_array = np.array([[predicted_value]])
    current_window = np.vstack([current_window[1:], predicted_value_as_array])

# Convert predictions back to original scale (if using a scaler)
predicted_prices = scaler.inverse_transform(np.array(noisy_predictions).reshape(-1, 1))

# Number of days to zoom in on
zoom_days = 185

# Create a timeline for the final `zoom_days` of the actual data and original predictions
days = np.arange(len(original_data))[-zoom_days:] # Last `zoom_days` days for the actual data
predicted_days = np.arange(len(original_data))[-zoom_days:] # Adjust for the predicted data's timeline
future_days = np.arange(len(original_data), len(original_data) + len(predicted_prices)) # Future timeline

# Noisy future predictions plot
plt.figure(figsize=(10, 6))
plt.plot(days, original_data[-zoom_days:], label='Actual Data', color='blue', linewidth=2)
plt.plot(predicted_days, predicted_data[-zoom_days:], label='Predicted Data', color='orange', linewidth=2)
plt.plot(future_days, predicted_prices, label='Future Predicted Data', color='#4CAF50', linewidth=2) # Green color for future predictions

# Add a vertical dotted line at the point where the future predictions start
plt.axvline(x=len(original_data) - 1, color='black', linestyle='--', label='Prediction Boundary')

# Add labels and title with bold fonts
plt.title(f'Future Stock Price Prediction for {ticker}', fontsize=18, weight='bold', pad=20)
plt.xlabel('Days', fontsize=14, weight='bold')
plt.ylabel('Stock Price', fontsize=14, weight='bold')

# Add grid and legend
plt.grid(color='grey', linestyle='--', linewidth=0.5, alpha=0.7)
plt.legend(loc='lower right', fontsize=12)

plt.savefig(f'{directory}/noisy_future.png')
plt.close()

```

testing the function

```

#test
pipeline('WM')

```

```

[*****100%*****] 1 of 1 completed
1/1 [=====] - 1s 1s/step
Epoch 1/100
100/100 [=====] - 8s 22ms/step - loss: 0.0206 - val_loss: 0.0394 - lr: 0.0010
Epoch 2/100
100/100 [=====] - 1s 11ms/step - loss: 0.0094 - val_loss: 0.0088 - lr: 0.0010

```

```

Epoch 3/100
100/100 [=====] - 1s 10ms/step - loss: 0.0069 - val_loss: 0.0054 - lr: 0.0010
Epoch 4/100
100/100 [=====] - 1s 11ms/step - loss: 0.0066 - val_loss: 0.0046 - lr: 0.0010
Epoch 5/100
100/100 [=====] - 1s 10ms/step - loss: 0.0077 - val_loss: 0.0047 - lr: 0.0010
Epoch 6/100
100/100 [=====] - 1s 11ms/step - loss: 0.0071 - val_loss: 0.0036 - lr: 0.0010
Epoch 7/100
100/100 [=====] - 1s 10ms/step - loss: 0.0066 - val_loss: 0.0035 - lr: 0.0010
Epoch 8/100
100/100 [=====] - 1s 12ms/step - loss: 0.0071 - val_loss: 0.0034 - lr: 0.0010
Epoch 9/100
100/100 [=====] - 1s 11ms/step - loss: 0.0070 - val_loss: 0.0036 - lr: 0.0010
Epoch 10/100
100/100 [=====] - 1s 11ms/step - loss: 0.0061 - val_loss: 0.0032 - lr: 0.0010
Epoch 11/100
100/100 [=====] - 1s 11ms/step - loss: 0.0063 - val_loss: 0.0035 - lr: 0.0010
Epoch 12/100
100/100 [=====] - 1s 11ms/step - loss: 0.0060 - val_loss: 0.0031 - lr: 0.0010
Epoch 13/100
100/100 [=====] - 1s 10ms/step - loss: 0.0057 - val_loss: 0.0029 - lr: 0.0010
Epoch 14/100
100/100 [=====] - 1s 11ms/step - loss: 0.0054 - val_loss: 0.0042 - lr: 0.0010
Epoch 15/100
100/100 [=====] - 1s 11ms/step - loss: 0.0054 - val_loss: 0.0046 - lr: 0.0010
Epoch 16/100
100/100 [=====] - 1s 10ms/step - loss: 0.0056 - val_loss: 0.0026 - lr: 0.0010
Epoch 17/100
100/100 [=====] - 1s 11ms/step - loss: 0.0051 - val_loss: 0.0023 - lr: 0.0010
Epoch 18/100
100/100 [=====] - 1s 10ms/step - loss: 0.0050 - val_loss: 0.0043 - lr: 0.0010
Epoch 19/100
100/100 [=====] - 1s 10ms/step - loss: 0.0049 - val_loss: 0.0023 - lr: 0.0010
Epoch 20/100
100/100 [=====] - 1s 11ms/step - loss: 0.0045 - val_loss: 0.0044 - lr: 0.0010
Epoch 21/100
100/100 [=====] - 1s 11ms/step - loss: 0.0051 - val_loss: 0.0029 - lr: 0.0010
Epoch 22/100
98/100 [=====>.] - ETA: 0s - loss: 0.0045
Epoch 22: ReduceLROnPlateau reducing learning rate to 0.00010000000474974513.
100/100 [=====] - 1s 11ms/step - loss: 0.0045 - val_loss: 0.0056 - lr: 0.0010
Epoch 23/100
100/100 [=====] - 1s 10ms/step - loss: 0.0050 - val_loss: 0.0023 - lr: 1.0000e-04
Epoch 24/100
100/100 [=====] - 1s 11ms/step - loss: 0.0036 - val_loss: 0.0022 - lr: 1.0000e-04
Epoch 25/100
100/100 [=====] - 1s 9ms/step - loss: 0.0039 - val_loss: 0.0021 - lr: 1.0000e-04
Epoch 26/100
100/100 [=====] - 1s 11ms/step - loss: 0.0036 - val_loss: 0.0020 - lr: 1.0000e-04
Epoch 27/100
100/100 [=====] - 1s 11ms/step - loss: 0.0033 - val loss: 0.0022 - lr: 1.0000e-04

```

IT WORKS!!!

✓ The Final Result

```

# List of 35 stocks with the highest market cap in respective sector
tickers = ['AAPL', 'MSFT', 'NVDA', 'AVGO', 'ORCL', 'ADBE', 'CRM', 'AMD', 'ACN', # 9 Stocks from the Technology Sector
           'BRK-B', 'JPM', 'V', 'MA', 'BAC', 'WFC', # 6 Stocks from the Financial Services Sector
           'LLY', 'UNH', 'JNJ', 'ABBV', # 4 Stocks from the Healthcare Sector
           'GE', 'CAT', 'RTX', 'UNP', # 4 Stocks from the Industrials Sector
           'AMZN', 'TSLA', 'HD', # 3 Stocks from the Consumer Cyclical Sector
           'GOOGL', 'META', 'NFLX', # 3 Stocks from the Communication Services Sector
           'WMT', 'PG', # 2 Stocks from the Consumer Defensive Sector
           'XOM', # 1 Stock from the Energy Sector
           'PLD', # 1 Stock from the Real Estate Sector
           'NEE', # 1 Stock from the Utilities Sector
           'LIN' # 1 Stock from the Basic Materials Sector
]

```

for ticker in tickers: